



An OSS API Layer for your Apache Cassandra™



STARGATE



Cédrick Lunven

Director Developer Relations



SDK



- Trainer
- Public Speaker
- Developers Support
- Developer Applications
- Developer Tooling



- Creator of ff4j (ff4j.org)
- Maintainer for 8 years+



- Happy developer for 14 years
- Spring Petclinic Reactive & Starters
- Implementing APIs for 8 years

DataStax Developers Advocates Team



Cedrick
Lunven

Aleksandr
Volochnev

Jack
Fryer

Kirsten
Hunter

Stefano
Lottini

David
Gilardi

Ryan
Welford

Rags
Srinivas

Sonia
Siganporia

R

S

Agenda (40 min)

01



Stargate Data gateway
What, Why and How

02 {REST}

Exploring Apis
Rest DML and DDL

03



Exploring Apis
Document oriented

04



Exploring Apis
GraphQL and Federation

05



Tooling
SDK, K8ssandra

05



What's NEXT ?
gRPC, CDC, sql...

Agenda (40 min)

01



Stargate Data gateway
What, Why and How

02 {REST}

Exploring Apis
Rest DML and DDL

03



Exploring Apis
Document oriented

04



Exploring Apis
GraphQL and Federation

05



Tooling
SDK, K8ssandra

05



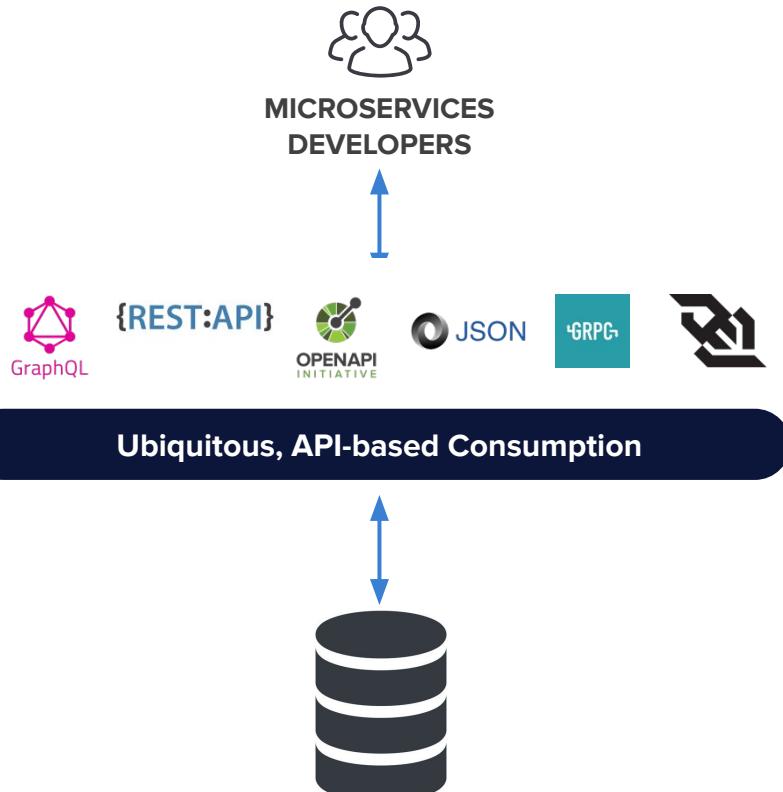
What's NEXT ?
gRPC, CDC, sql...

Data Gateway



Developers want the option to use modern APIs and development gateways.

Cassandra is a database designed for the new standard and the associated APIs that facilitate the first choice of developers.



Stargate Overview

An open source API framework for data

Stargate makes it easy to use a database for any application workload by adding plugin support for new APIs, data types, and access methods

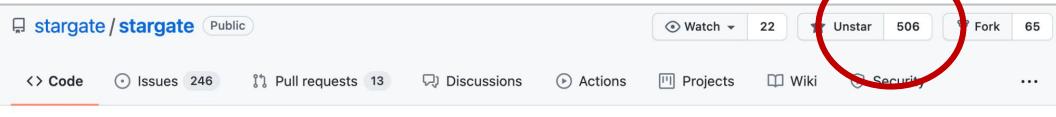


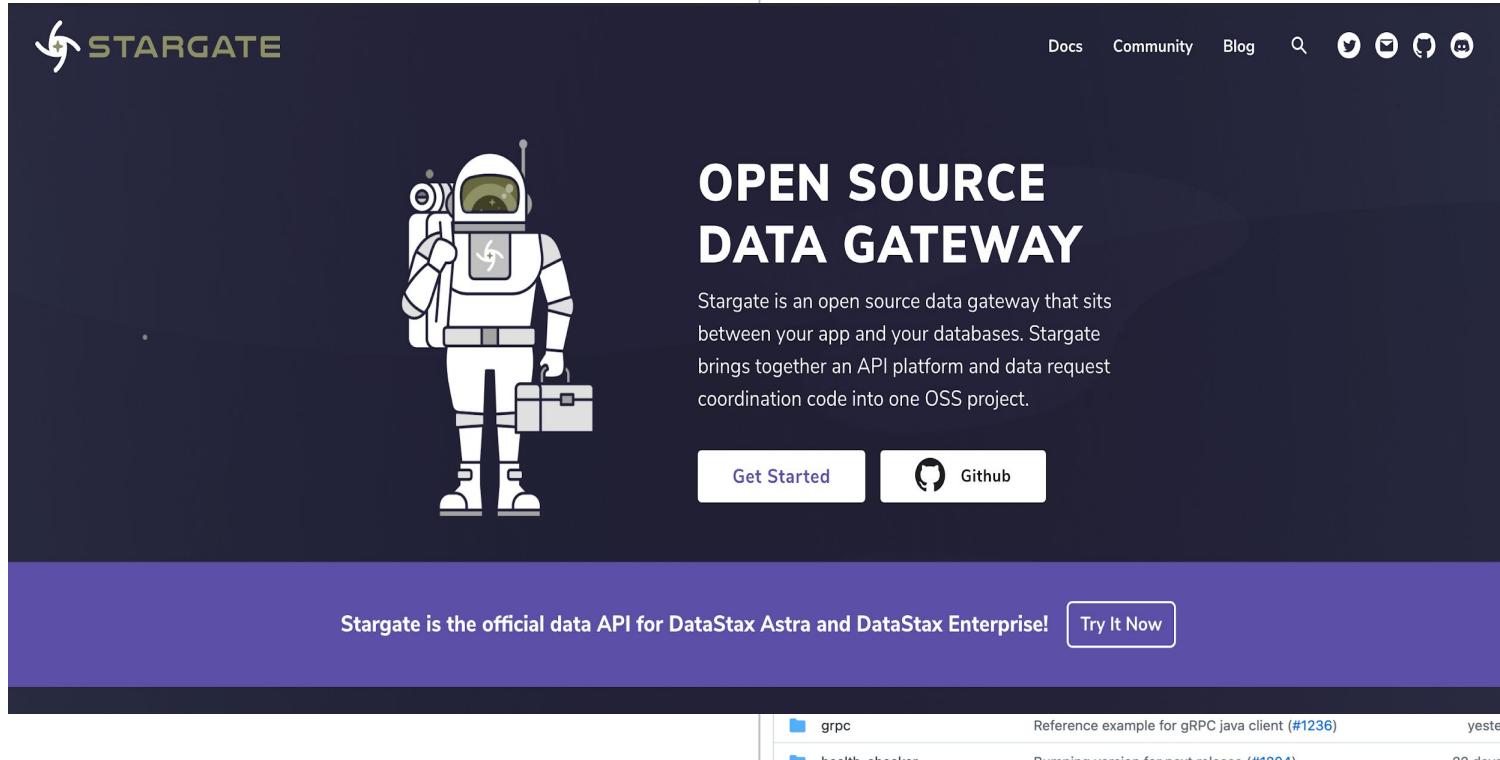
MICROSERVICES
DEVELOPERS



STARGATE







STARGATE

OPEN SOURCE DATA GATEWAY

Stargate is an open source data gateway that sits between your app and your databases. Stargate brings together an API platform and data request coordination code into one OSS project.

Get Started  Github

Stargate is the official data API for DataStax Astra and DataStax Enterprise! 

Reference example for gRPC java client (#1236) 

About
An open source data gateway
 stargate.io
java graphql rest
cassandra cql

Readme
Apache-2.0 License

Releases 62
 Release v1.0.32 (Latest)
25 days ago
+ 61 releases

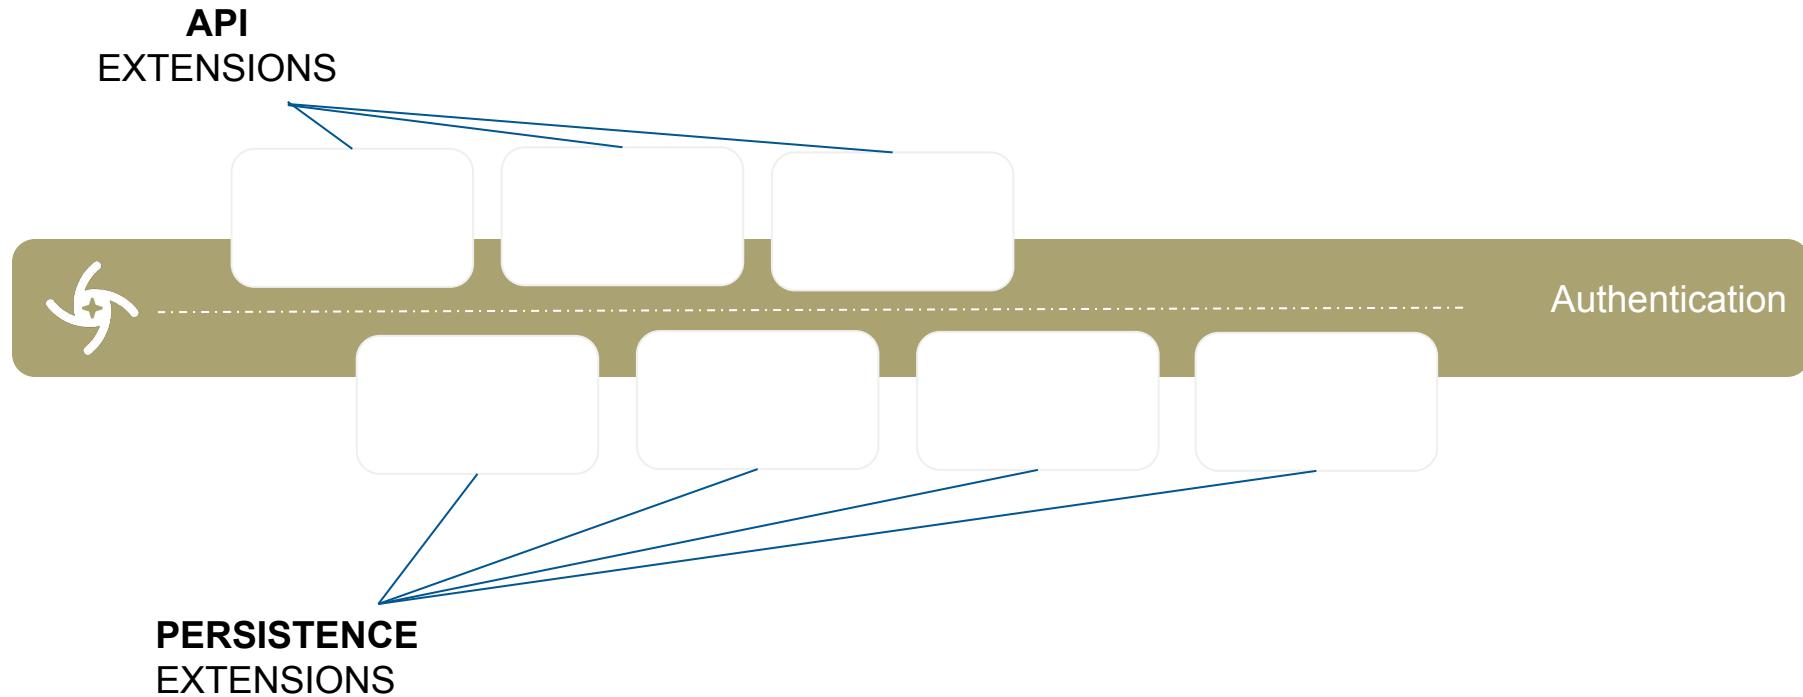
Packages 13



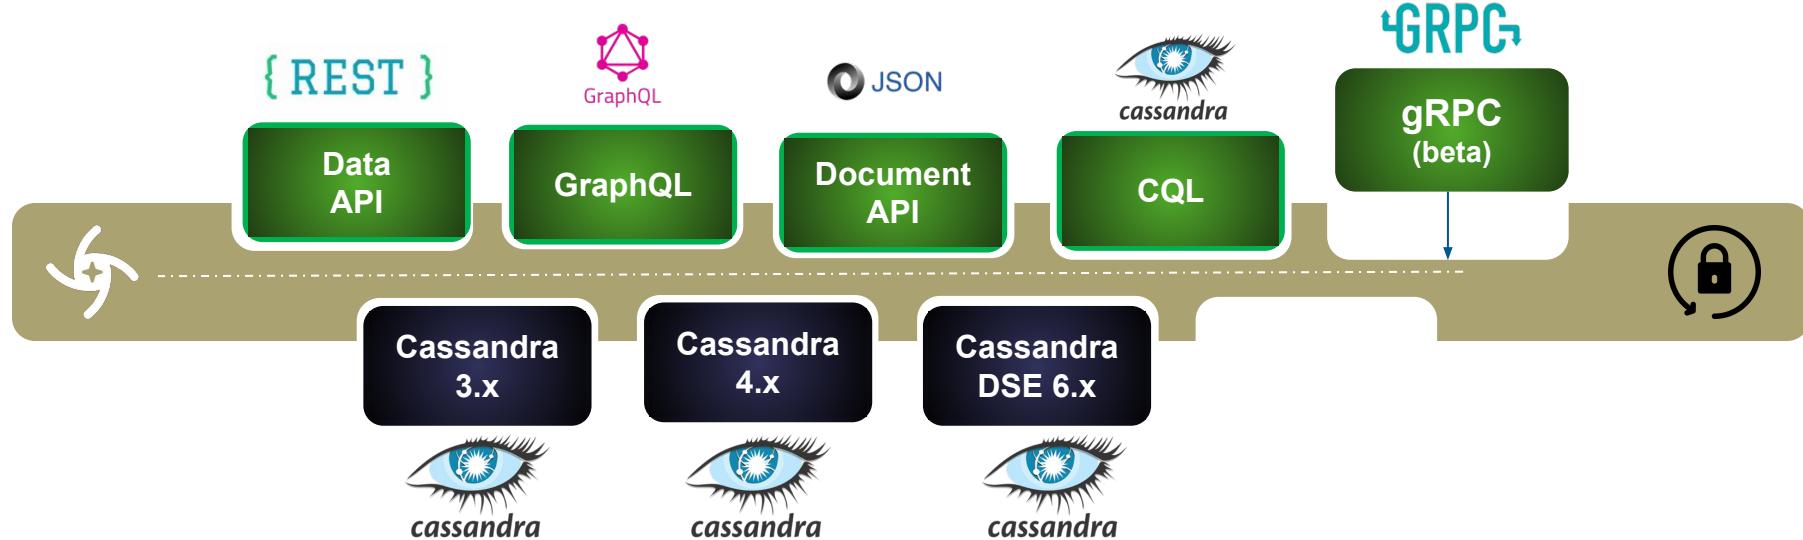
+ 10 packages

Contributors 30 

API and Persistence Extensions



API Extensions and Persistence Extensions



9042

8090

8082

8080

8084

CQL

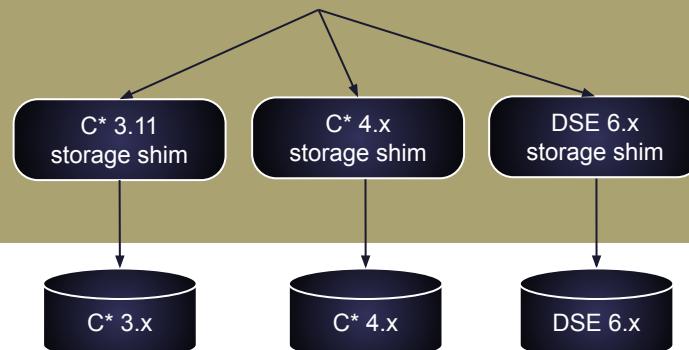
GRPC

Data + Document APIs

GraphQL APIs



STARGATE



9042

8090

8081

8082

8080

8084

CQL

GRPC

Auth

Data + Document APIs

GraphQL APIs

Authentication Service

Rate Limiting Service

Metrics Services



STARGATE



Persistence Service

C* 3.11
storage shimC* 4.x
storage shimDSE 6.x
storage shimOSGI
Service Registry

APACHECON

DataStax Developers

9042

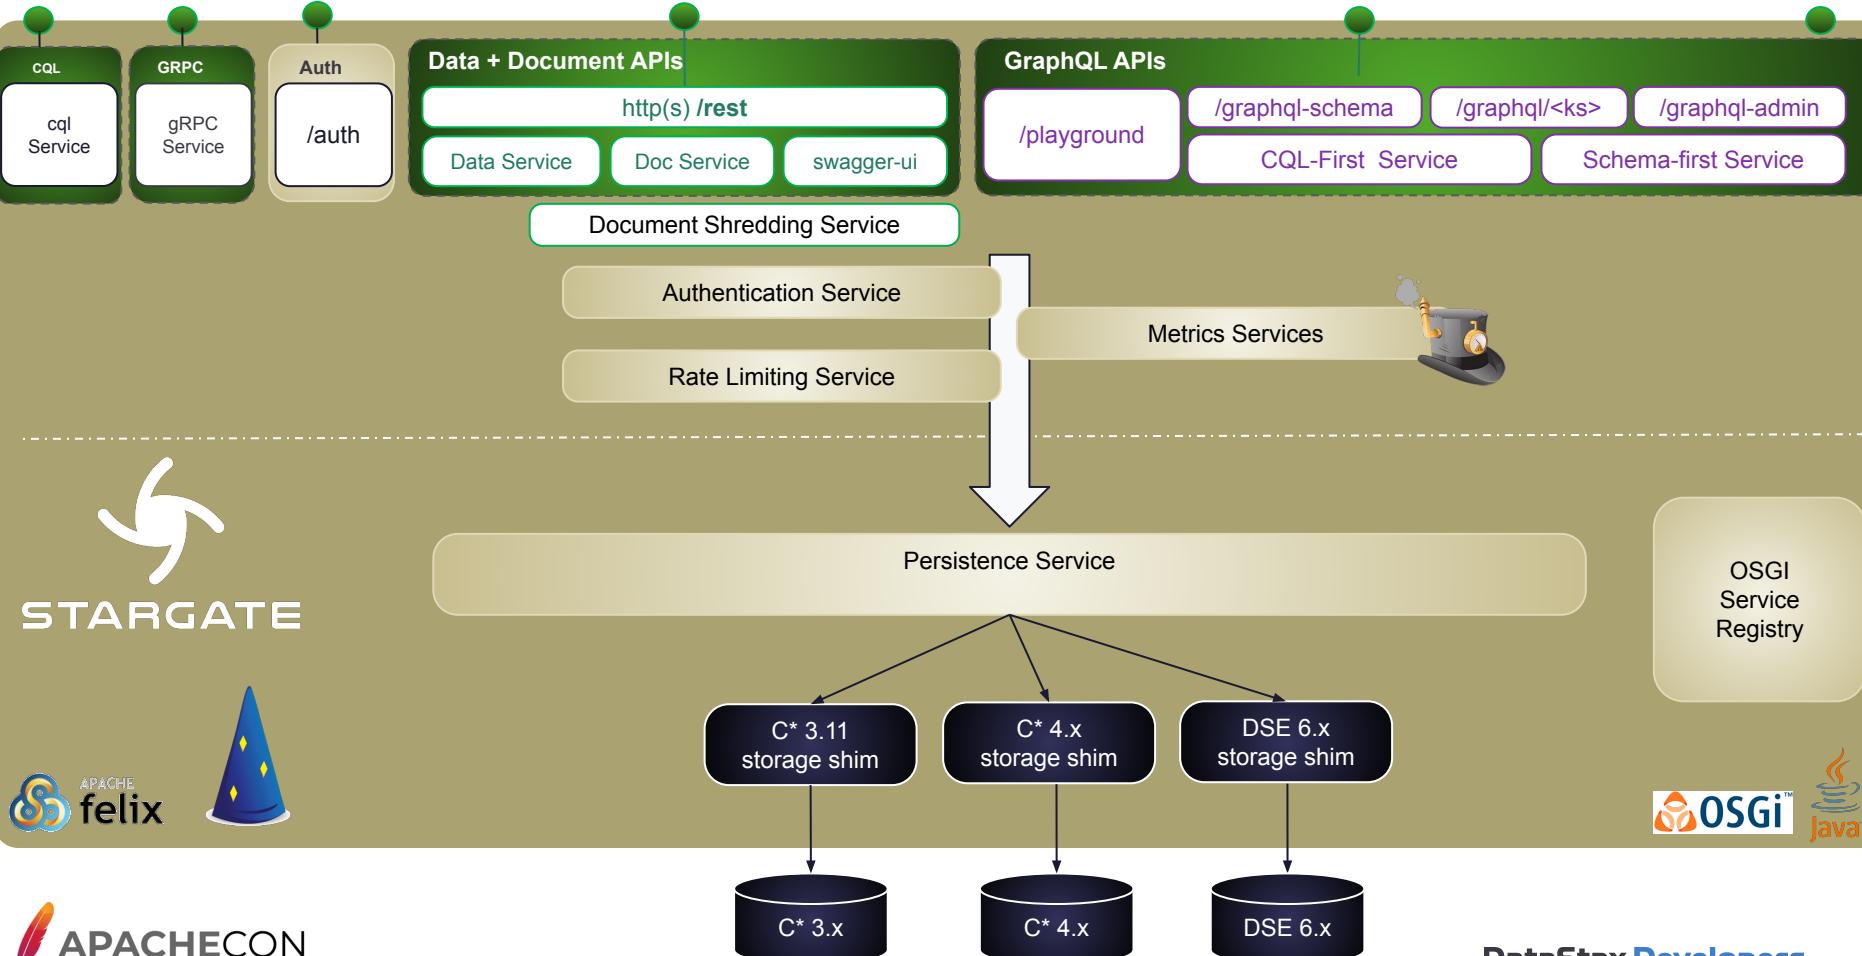
8090

8081

8082

8080

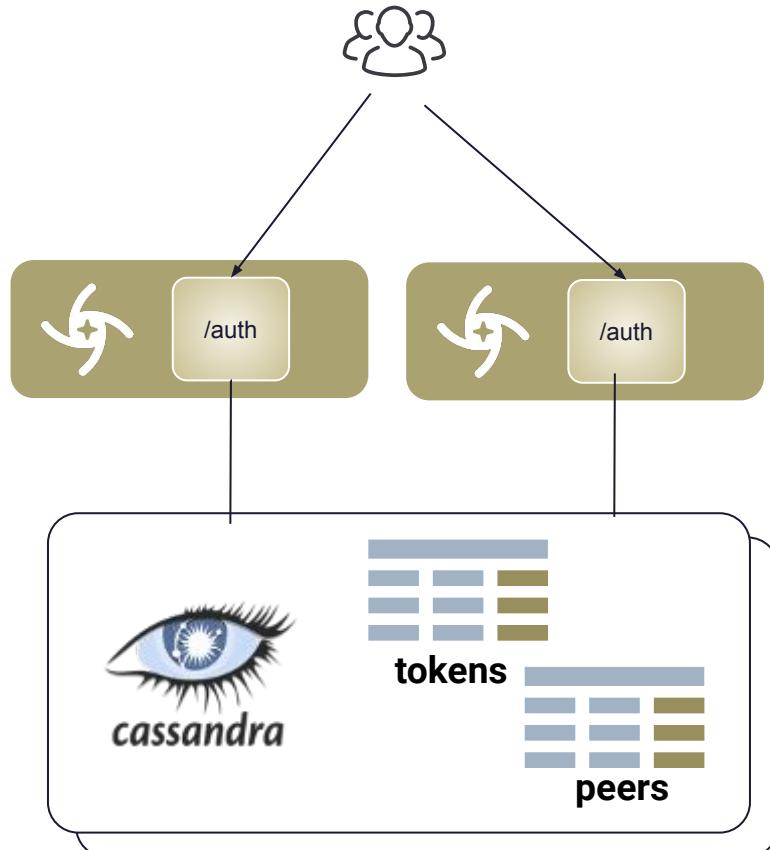
8084



Authentication / Authorization

- Use Auth and get tokens (user+pwd)
- Tokens shared across instances
- Leverage on Cassandra roles (RBAC)
- -Dstargate.auth_tokenttl =X

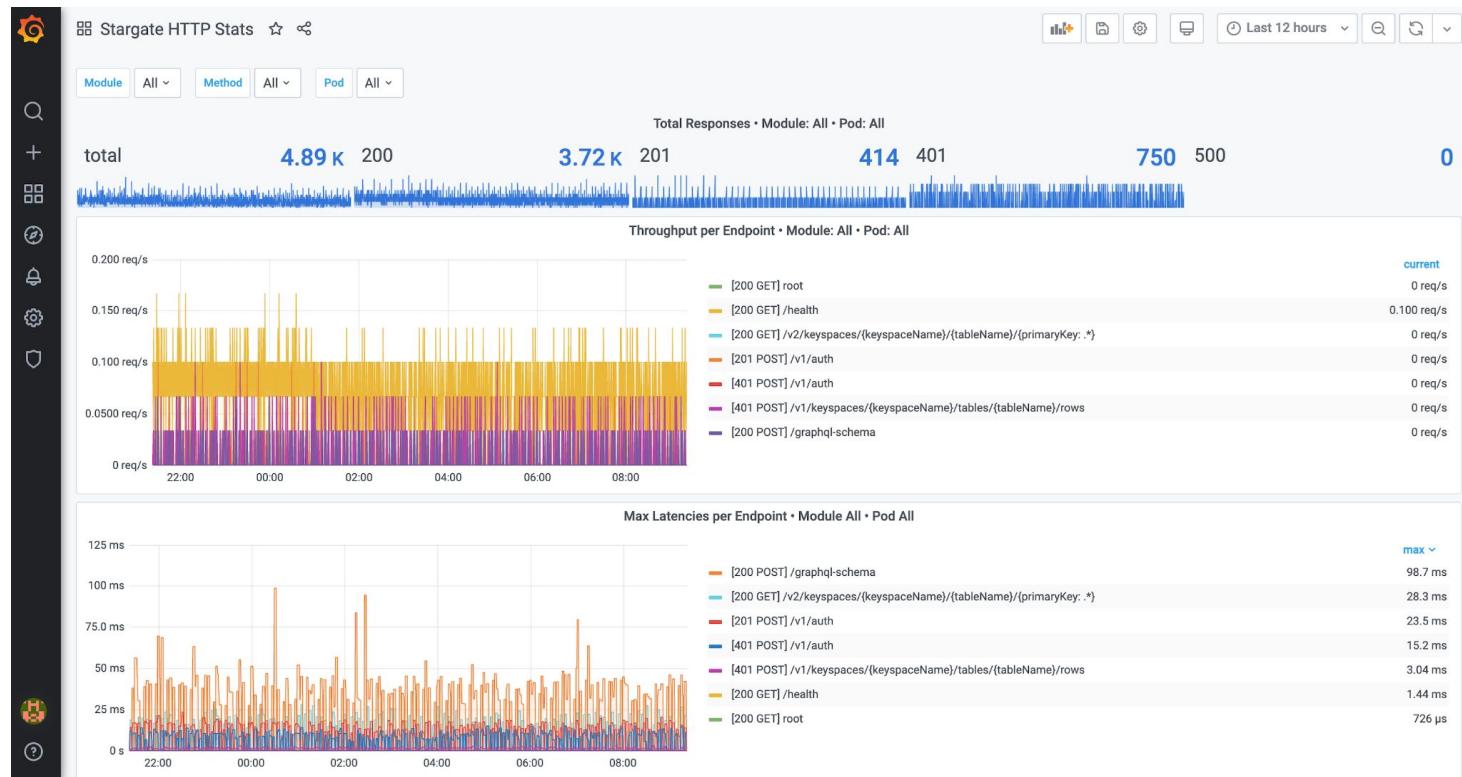
```
curl -L -X POST 'http://localhost:8081/v1/auth' \
-H 'Content-Type: application/json' \
--data-raw '{
    "username": "cassandra",
    "password": "cassandra"
}'
```



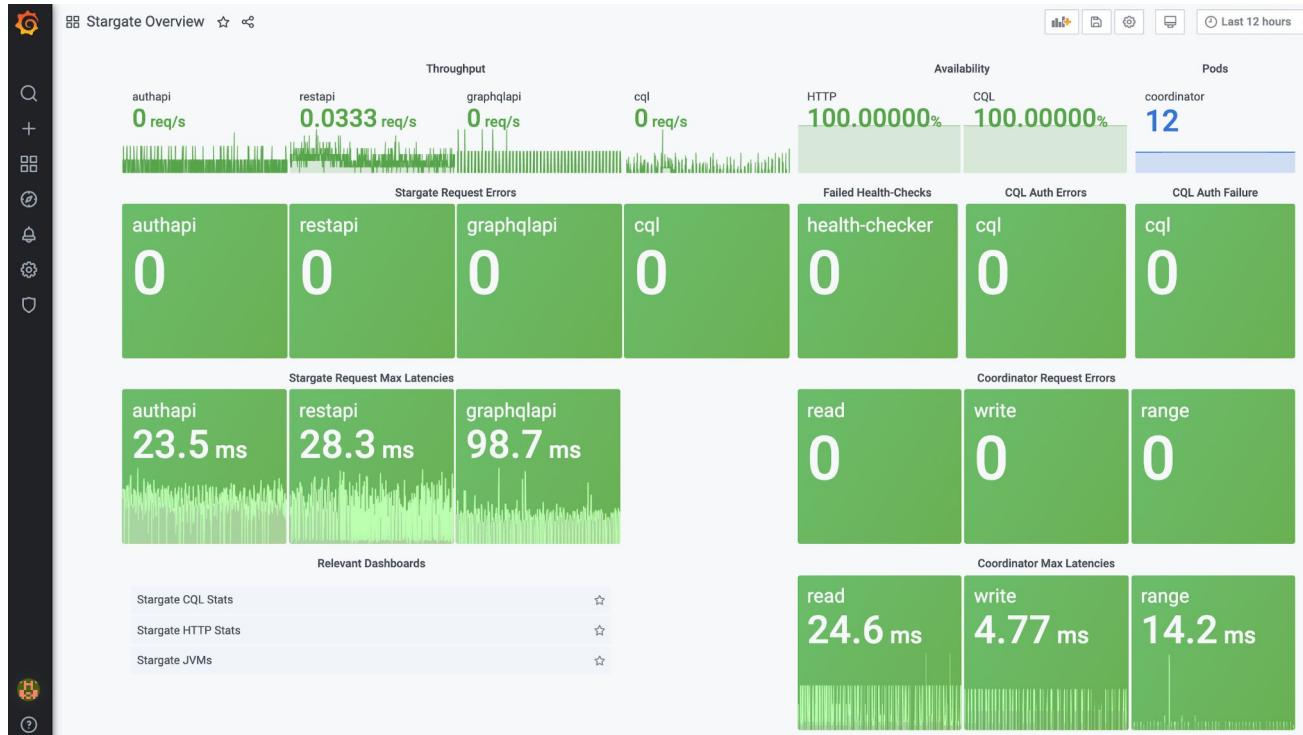
Monitoring (JVM)



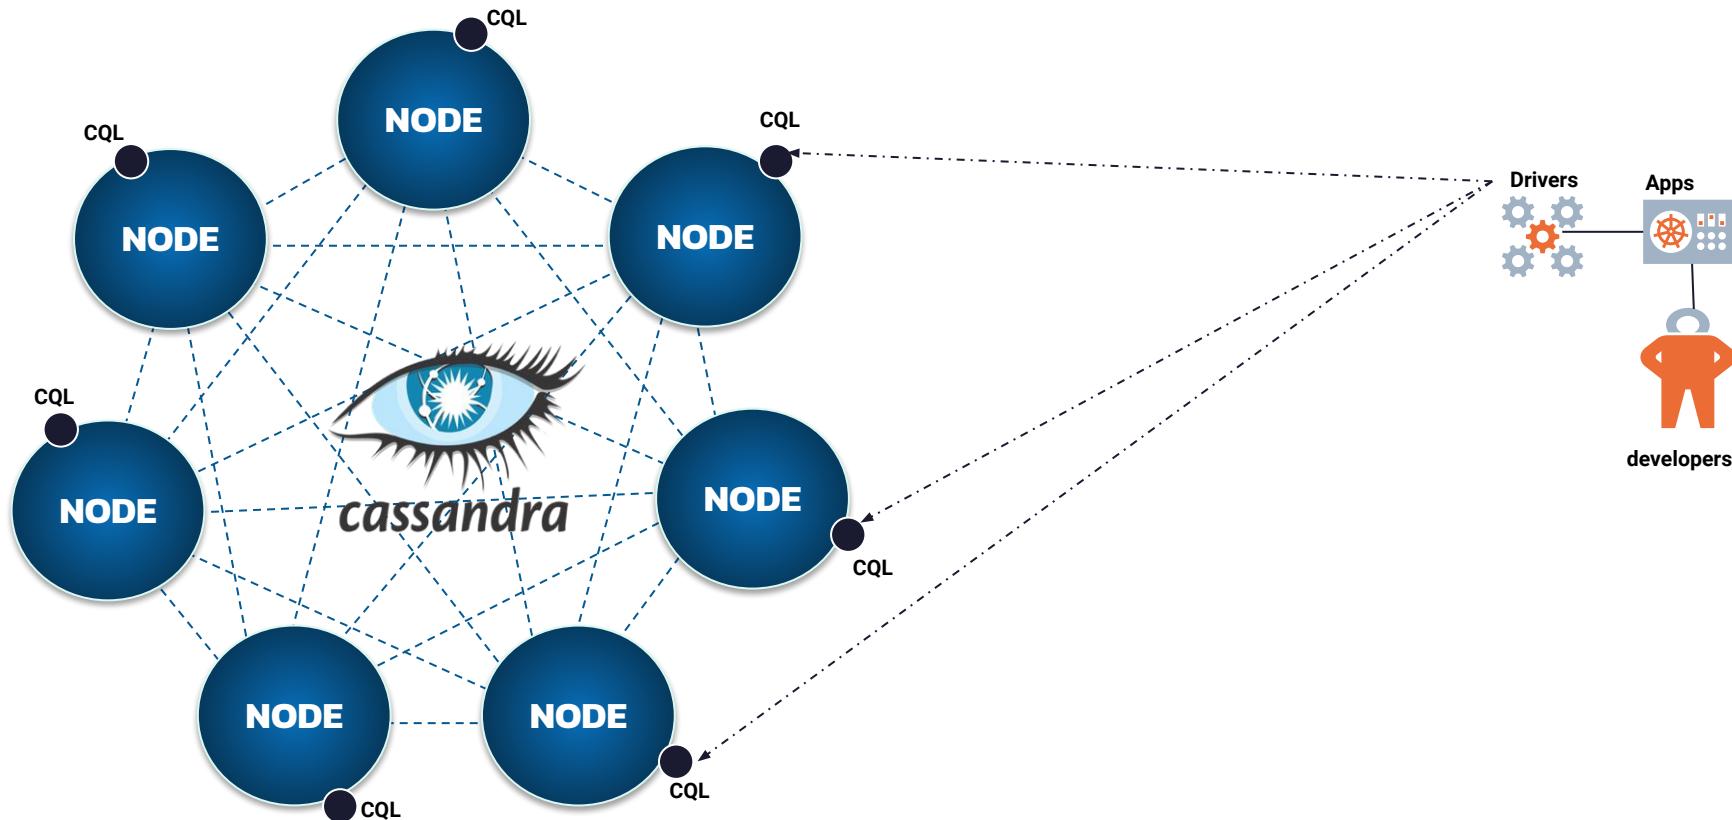
Monitoring (HTTP Traffic)



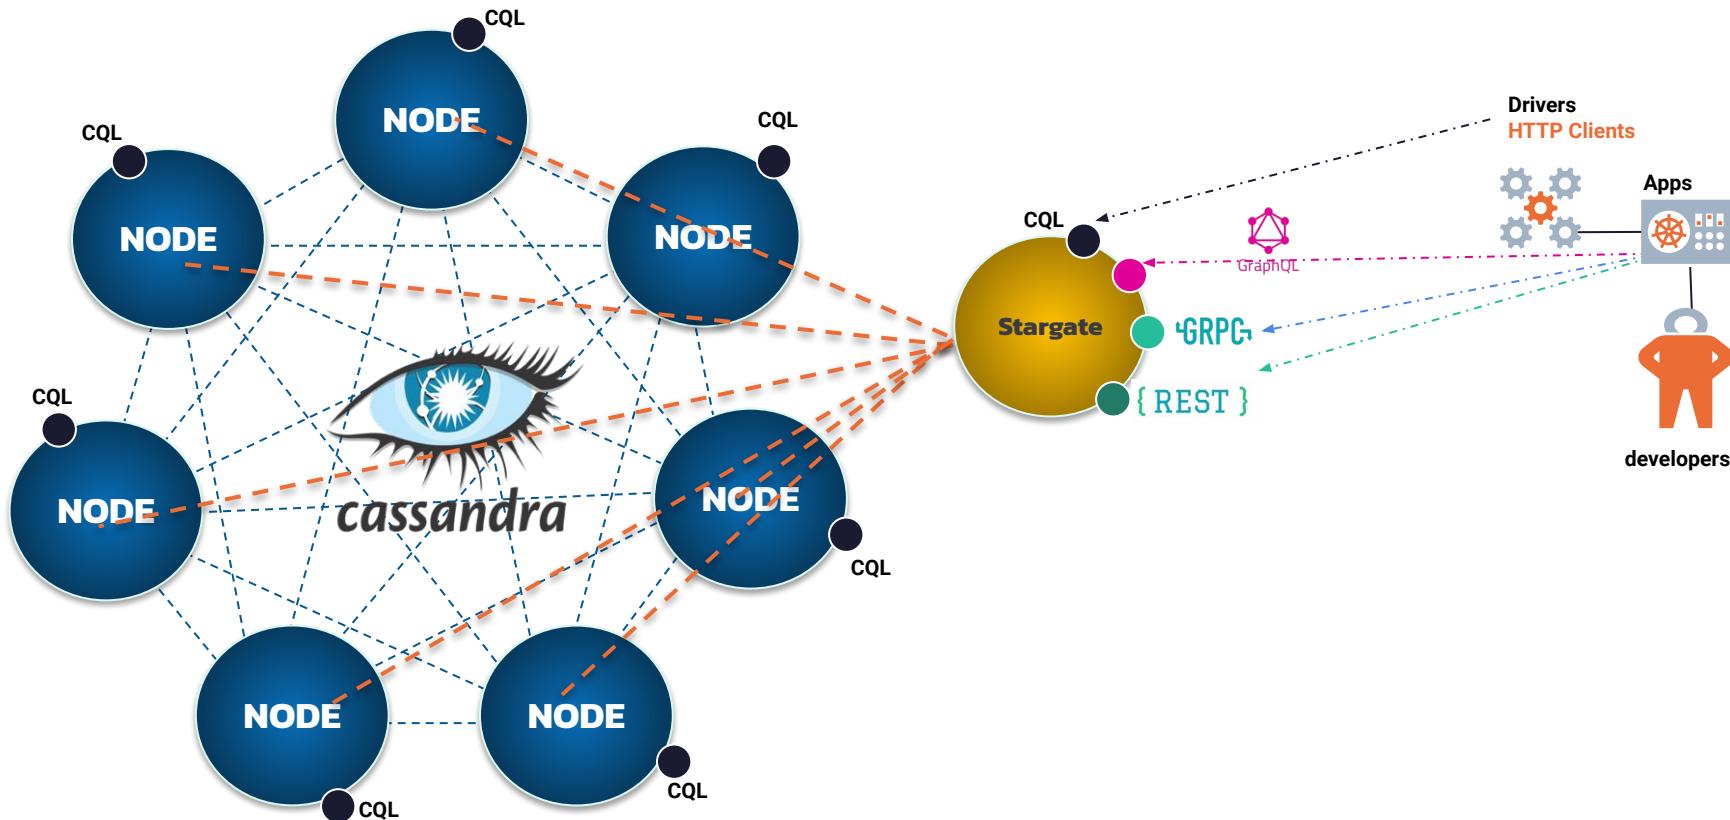
Monitoring (Services)



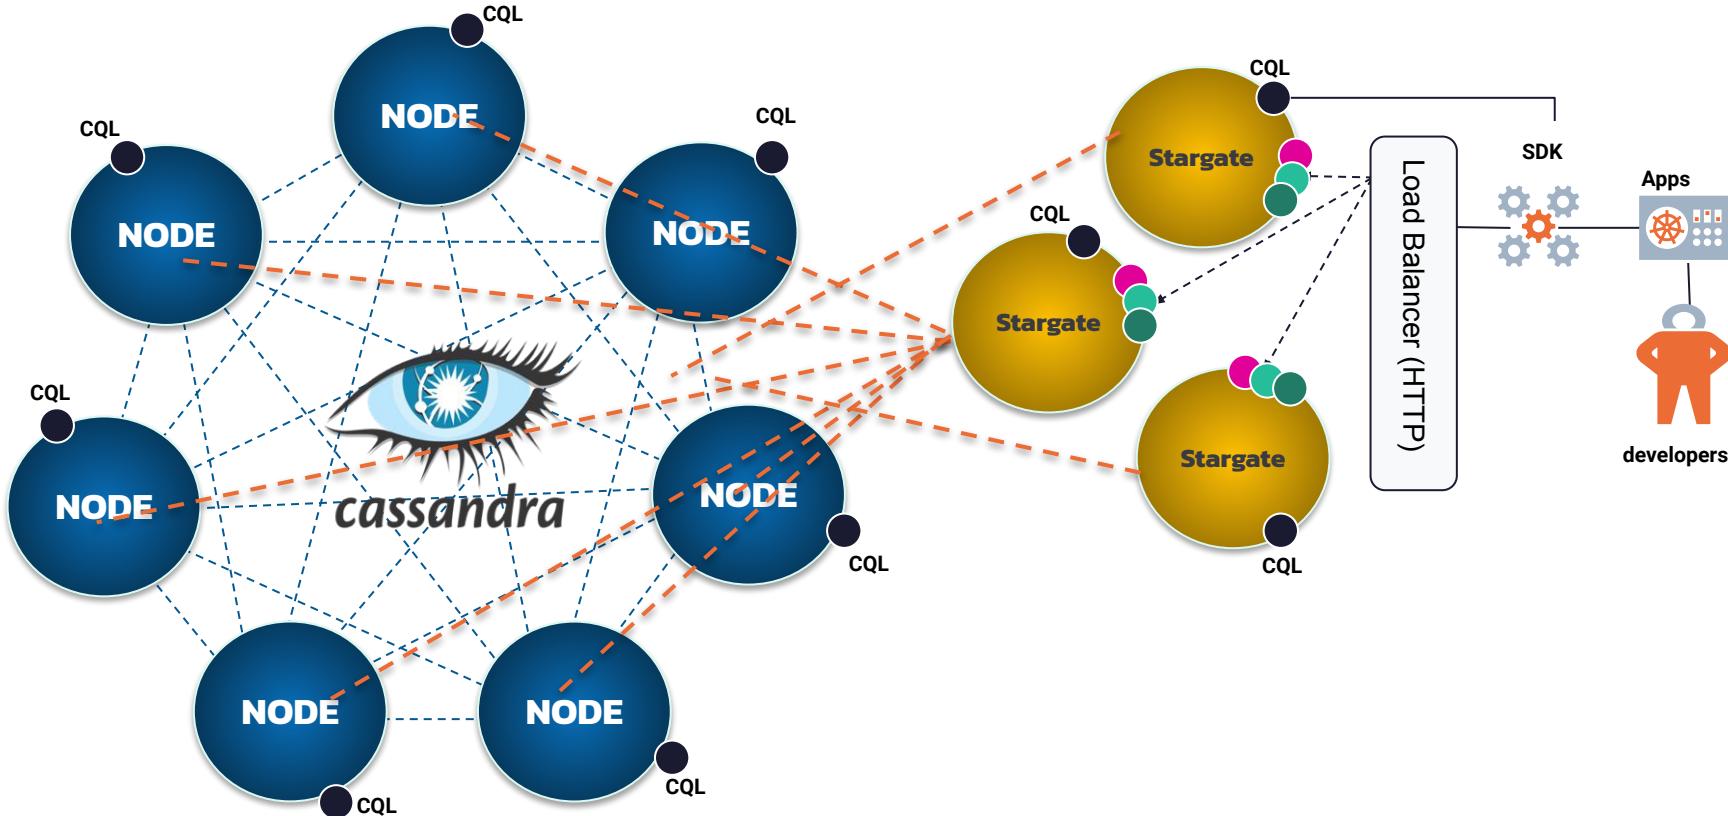
Connecting to your cluster (Before)



Connecting to your cluster (with Stargate)



Connecting to your cluster (with Stargates)



Agenda (40 min)

01



Stargate Data gateway
What, Why and How

02 {REST}

Exploring Apis
Rest DML and DDL

03



Exploring Apis
Document oriented

04



Exploring Apis
GraphQL and Federation

05



Tooling
SDK, K8ssandra

05



What's NEXT ?
gRPC, CDC, sql...

Stargate “REST” Api(s)

- Expose existing CQL schema as a RESTful endpoint, or create new CQL schema
- Most familiar API style for most teams
- Swagger integration

schemas

GET	/v1/keyspaces/{keyspaceName}/tables/{tableName}/columns	Retrieve all columns
POST	/v1/keyspaces/{keyspaceName}/tables/{tableName}/columns	Add a column
GET	/v1/keyspaces/{keyspaceName}/tables/{tableName}/columns/{columnName}	Retrieve a column
PUT	/v1/keyspaces/{keyspaceName}/tables/{tableName}/columns/{columnName}	Update a column
DELETE	/v1/keyspaces/{keyspaceName}/tables/{tableName}/columns/{columnName}	Delete a column
GET	/v1/keyspaces	Return all keyspaces
GET	/v1/keyspaces/{keyspaceName}/tables	Return all tables
POST	/v1/keyspaces/{keyspaceName}/tables	Add a table
GET	/v1/keyspaces/{keyspaceName}/tables/{tableName}	Return a table
DELETE	/v1/keyspaces/{keyspaceName}/tables/{tableName}	Delete a table
GET	/v2/schemas/keyspaces/{keyspaceName}/tables/{tableName}/columns	Get all columns

Quick start: https://stargate.io/docs/stargate/1.0/quickstart/quick_start-rest.html

REST Api(s) exposed DDL and best practices DML

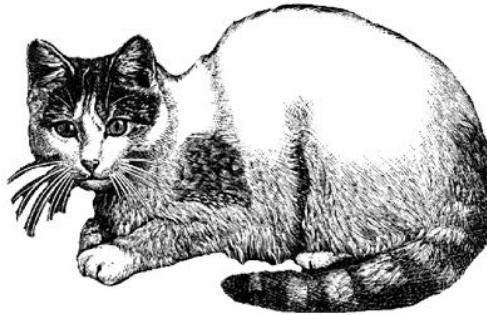
GET	/v2/schemas/keyspaces/{keyspaceName}/tables/{tableName}/columns	Get all columns
POST	/v2/schemas/keyspaces/{keyspaceName}/tables/{tableName}/columns	Create a column
GET	/v2/schemas/keyspaces/{keyspaceName}/tables/{tableName}/columns/{columnName}	Get a column
PUT	/v2/schemas/keyspaces/{keyspaceName}/tables/{tableName}/columns/{columnName}	Update a column
DELETE	/v2/schemas/keyspaces/{keyspaceName}/tables/{tableName}/columns/{columnName}	Delete a column
DELETE	/v2/schemas/keyspaces/{keyspaceName}/tables/{tableName}/indexes/{indexName}	Drop an index from keyspace
GET	/v2/schemas/keyspaces/{keyspaceName}/tables/{tableName}/indexes	Get all indexes for a given table
POST	/v2/schemas/keyspaces/{keyspaceName}/tables/{tableName}/indexes	Add an index to a table's column
GET	/v2/schemas/keyspaces/{keyspaceName}	Get a keyspace
DELETE	/v2/schemas/keyspaces/{keyspaceName}	Delete a keyspace
GET	/v2/schemas/keyspaces	Get all keyspaces
POST	/v2/schemas/keyspaces	Create a keyspace
GET	/v2/schemas/keyspaces/{keyspaceName}/tables/{tableName}	Get a table
PUT	/v2/schemas/keyspaces/{keyspaceName}/tables/{tableName}	Replace a table definition

DDL (Schema)

GET	/v2/keyspaces/{keyspaceName}/{tableName}/rows	Retrieve all rows
GET	/v2/keyspaces/{keyspaceName}/{tableName}	Search a table
POST	/v2/keyspaces/{keyspaceName}/{tableName}	Add row
GET	/v2/keyspaces/{keyspaceName}/{tableName}/{primaryKey}	Get row(s)
PUT	/v2/keyspaces/{keyspaceName}/{tableName}/{primaryKey}	Replace row(s)
DELETE	/v2/keyspaces/{keyspaceName}/{tableName}/{primaryKey}	Delete row(s)
PATCH	/v2/keyspaces/{keyspaceName}/{tableName}/{primaryKey}	Update part of a row(s)

Demo

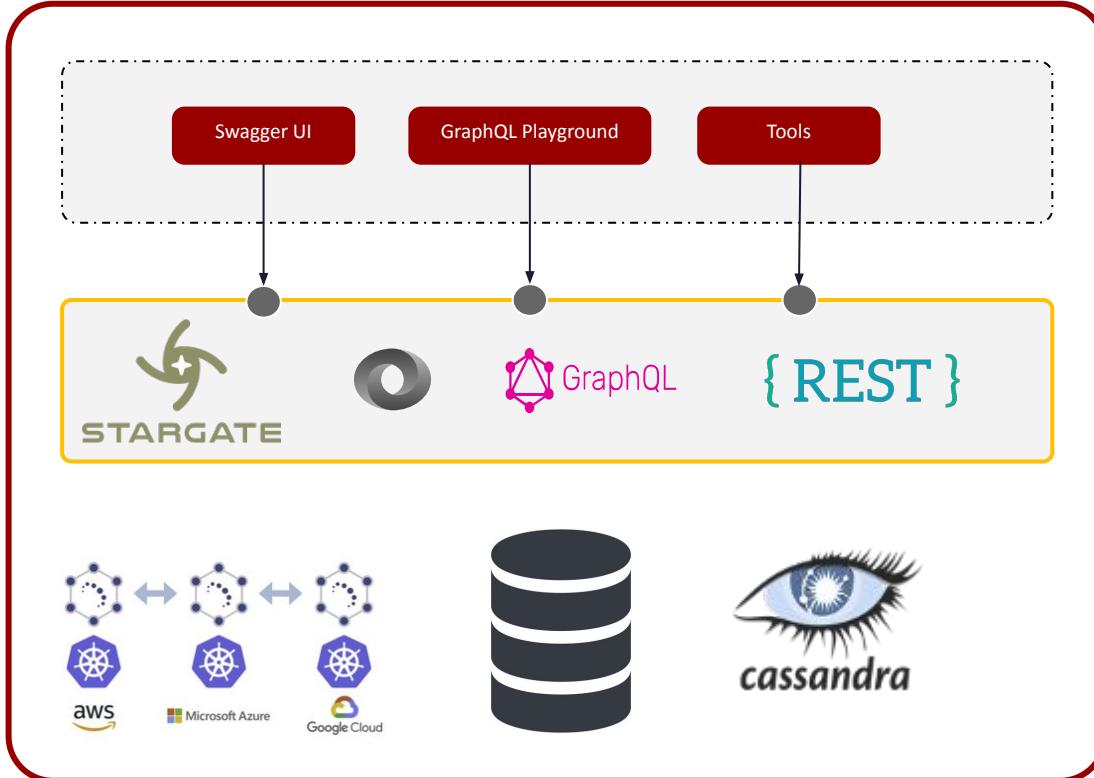
Stargate, the Developers Guide



More REST API
Less ALLOW FILTERING

O RLY[?]

Stackoverflow Expert



Agenda (40 min)

01



Stargate Data gateway
What, Why and How

02 {REST}

Exploring Apis
Rest DML and DDL

03



Exploring Apis
Document oriented

04



Exploring Apis
GraphQL and Federation

05



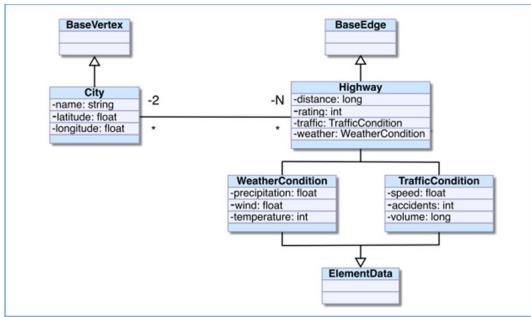
Tooling
SDK, K8ssandra

05



What's NEXT ?
gRPC, CDC, sql...

Working with JSON and Cassandra

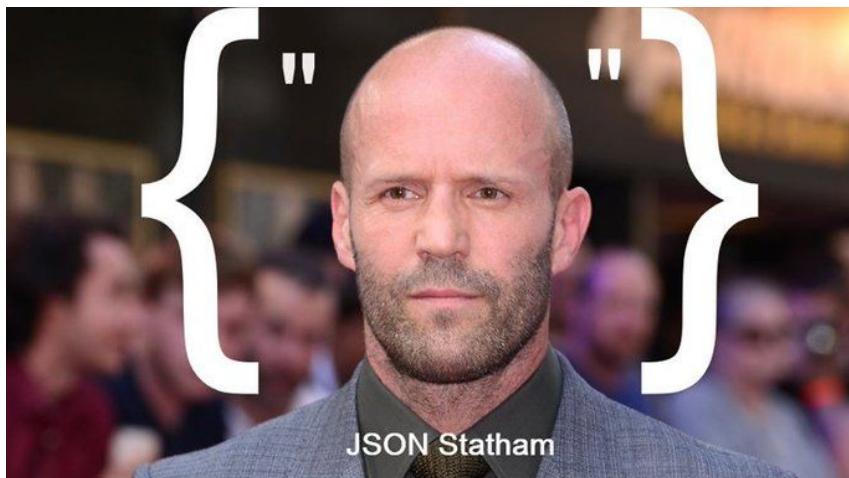


- Cassandra already handles both **JSON** and **nested objects**
 - INSERT JSON, SELECT JSON
 - Set<>, List<>, Map<>, and User Defined type (UDT) even nested
- But...
 - Strongly coupled with a SCHEMA VALIDATION
 - Dangerous with tombstones on updates

```
select json title,url,tags from videos;
```

```
INSERT INTO videos JSON '{
  "videoid": "e466f561-4ea4-4eb7-8dcc-126e0fbfd578",
  "email": "clunven@example.com",
  "title": "A JSON videos",
  "upload": "2020-02-26 15:09:22 +00:00",
  "url": "http://google.fr",
  "frames": [1,2,3,4],
  "tags": [ "cassandra", "accelerate", "2020"]
}';
```

“Schemaless”



“What rules ?”

- You want to insert and retrieve any JSON documents efficiently
- Allow “schemaless” (Validation less)
- Write to a single document is a single batch of statements
- Read from a single document is a single SELECT statement.
- Limit Tombstones with range deletes

Stargate Document API

- Schemaless!
(validationless)
- Add any JSON document
to a collection
- Uses “Document
shredding” approach

The screenshot shows the Stargate Document API documentation generated by Swagger. At the top, there's a header with the Stargate logo, the text "Supported by SMARTBEAR", and a "Explore" button. Below the header, the title "documents" is displayed. A list of API endpoints is shown, each with a method (e.g., GET, POST, DELETE, PUT, PATCH), a URL template, and a brief description. The endpoints are color-coded: blue for GET, green for POST, red for DELETE, orange for PUT, and teal for PATCH. The descriptions provide details like "List collections in namespace" or "Create a new empty collection in a namespace".

Method	URL	Description
GET	/v2/namespaces/{namespace-id}/collections	List collections in namespace
POST	/v2/namespaces/{namespace-id}/collections	Create a new empty collection in a namespace
GET	/v2/namespaces/{namespace-id}/collections/{collection-id}	Search documents in a collection
POST	/v2/namespaces/{namespace-id}/collections/{collection-id}	Create a new document
DELETE	/v2/namespaces/{namespace-id}/collections/{collection-id}	Delete a collection in a namespace
POST	/v2/namespaces/{namespace-id}/collections/{collection-id}/upgrade	Upgrade a collection in a namespace
POST	/v2/namespaces/{namespace-id}/collections/{collection-id}/batch	Write multiple documents in one request
GET	/v2/namespaces/{namespace-id}/collections/{collection-id}/{document-id}	Get a document
PUT	/v2/namespaces/{namespace-id}/collections/{collection-id}/{document-id}	Create or update a document with the provided document-id
DELETE	/v2/namespaces/{namespace-id}/collections/{collection-id}/{document-id}	Delete a document
PATCH	/v2/namespaces/{namespace-id}/collections/{collection-id}/{document-id}	Update data at the root of a document
GET	/v2/namespaces/{namespace-id}/collections/{collection-id}/{document-id}/{document-path}	Get a path in a document
PUT	/v2/namespaces/{namespace-id}/collections/{collection-id}/{document-id}/{document-path}	Replace data at a path in a document
DELETE	/v2/namespaces/{namespace-id}/collections/{collection-id}/{document-id}/{document-path}	Delete a path in a document
PATCH	/v2/namespaces/{namespace-id}/collections/{collection-id}/{document-id}/{document-path}	Update data at a path in a document
GET	/v2/namespaces/{namespace-id}/collections/{collection-id}/json-schema	Get a JSON schema from a collection

Quick start: https://stargate.io/docs/stargate/1.0/quickstart/quick_start-document.html

Document Shredding (1 / 3)

```
create table <name> (
    key text,
    p0 text,
    ... p[N] text,
    bool_value boolean,
    txt_value text, d
    bl_value double, leaf text
)
```

Document Shredding (2 / 3)

```
{"a": { "b": 1 }, "c": 2}
```

The document would be “shredded” into rows looking like this:

key	p0	p1	dbl_value
x	a	b	1
x	c	null	2

Document Shredding (3 / 3)

For data with an array, such as:

```
{"a": { "b": 1 }, "c": [{"d": 2}]}
```

there would be two rows, like so:

key	p0	p1	p2	dbl_value
x	a	b	null	1
x	c	[0]	d	2

Demo

Because we are sick of Mongo



Document Oriented
Apis like a Champ`

O RLY?

A frontend guy

Agenda (40 min)

01



Stargate Data gateway
What, Why and How

02 {REST}

Exploring Apis
Rest DML and DDL

03



Exploring Apis
Document oriented

04



Exploring Apis
GraphQL and Federation

05



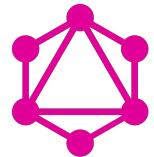
Tooling
SDK, K8ssandra

05



What's NEXT ?
gRPC, CDC, sql...

History



GraphQL

- 2012 Created by Facebook
Used internally for mobile apps
- 2015 Facebook give talk at ReactJs Conf
- 2015 Facebook open sourced GraphQL
- 2016 Github announced move to GQL

Definition

GraphQL is an application programming interface (API) query language and server-side runtime that prioritises giving customers precisely the data they request.

Why ?



GraphQL

Describe your data

```
type Project {  
  name: String  
  tagline: String  
  contributors: [User]  
}
```

Ask for what you want

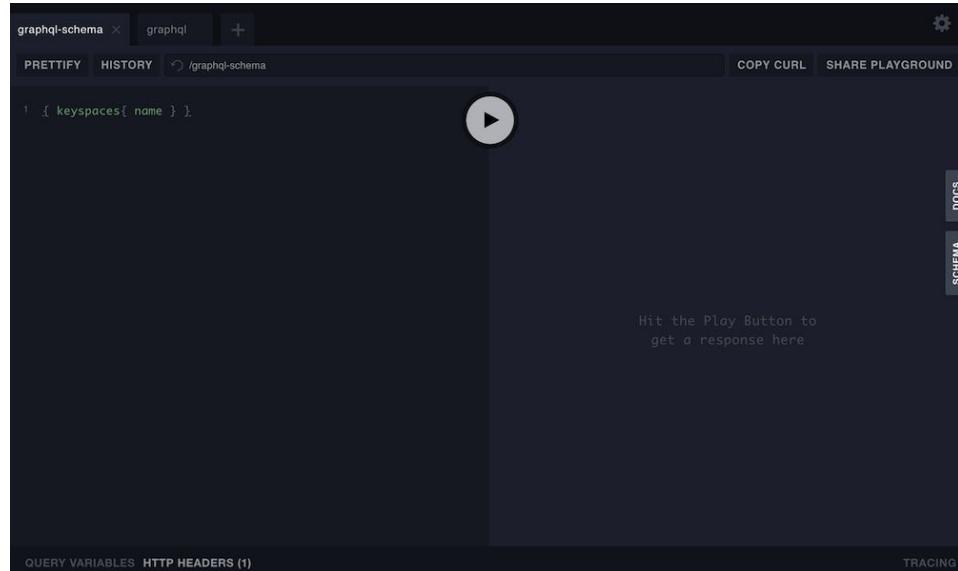
```
{  
  project(name: "GraphQL") {  
    tagline  
  }  
}
```

Get predictable results

```
{  
  "project": {  
    "tagline": "A query language for APIs"  
  }  
}
```

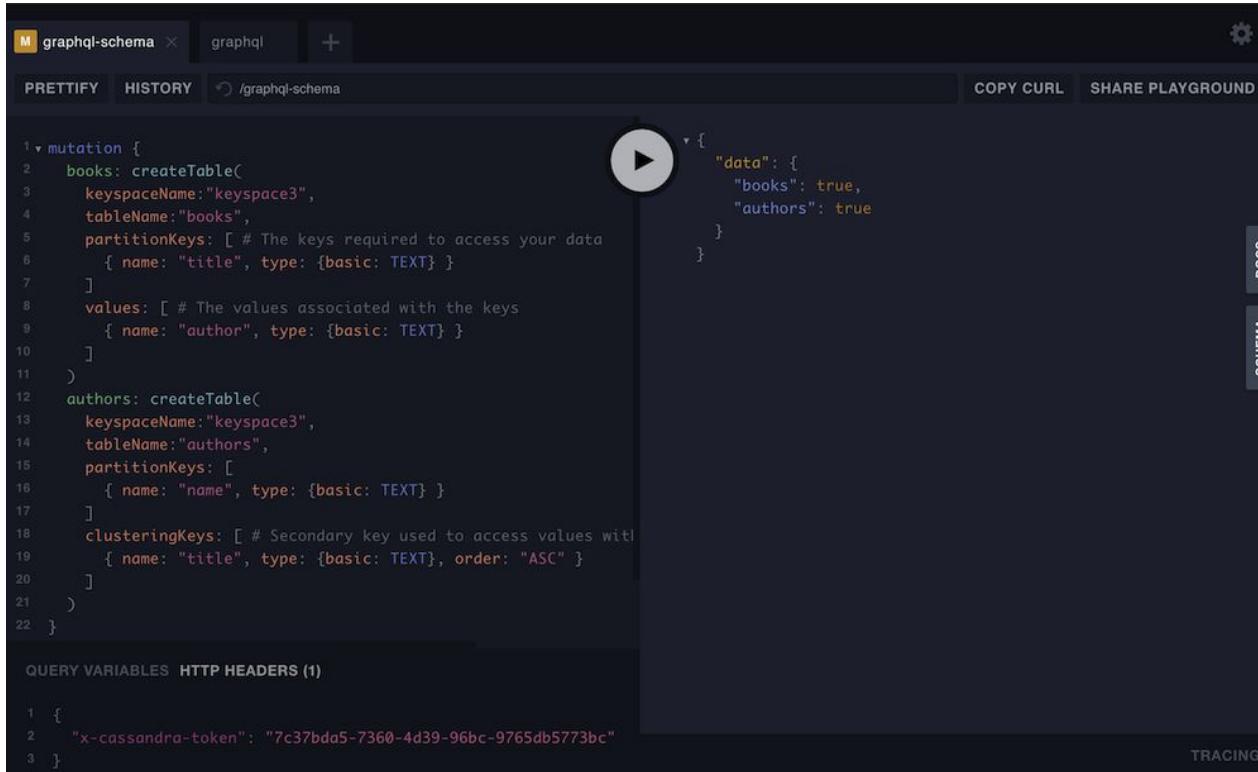
Stargate GraphQL API

- Expose existing CQL schema as a GraphQL endpoint
- Allows more fine-grained control of what fields returned
- In progress: federation across multiple services to join data
- Easy experimentation in GraphQL playground



Quick start: https://stargate.io/docs/stargate/1.0/quickstart/quick_start-graphql.html

GraphQL “CQL First”



The screenshot shows the GraphQL playground interface with a dark theme. The top navigation bar includes tabs for "graphql-schema" (selected), "graphql", and a "+" button. Below the tabs are buttons for "PRETTYIFY", "HISTORY", and a link to "/graphql-schema". On the right side, there are buttons for "COPY CURL", "SHARE PLAYGROUND", "SETTINGS", "DOCS" (selected), and "SCHEMA". A large play button is centered in the middle of the screen.

```
1 mutation {
2   books: createTable(
3     keyspaceName:"keyspace3",
4     tableName:"books",
5     partitionKeys: [ # The keys required to access your data
6       { name: "title", type: {basic: TEXT} }
7     ]
8     values: [ # The values associated with the keys
9       { name: "author", type: {basic: TEXT} }
10    ]
11  )
12   authors: createTable(
13     keyspaceName:"keyspace3",
14     tableName:"authors",
15     partitionKeys: [
16       { name: "name", type: {basic: TEXT} }
17     ]
18     clusteringKeys: [ # Secondary key used to access values with
19       { name: "title", type: {basic: TEXT}, order: "ASC" }
20     ]
21   )
22 }
```

Below the code editor, there are sections for "QUERY VARIABLES" and "HTTP HEADERS (1)".

QUERY VARIABLES

```
1 {  
2   "x-cassandra-token": "7c37bda5-7360-4d39-96bc-9765db5773bc"  
3 }
```

HTTP HEADERS (1)

```
1 {  
2   "x-cassandra-token": "7c37bda5-7360-4d39-96bc-9765db5773bc"  
3 }
```

On the bottom right, there is a "TRACING" button.

GraphQL “CQL First”

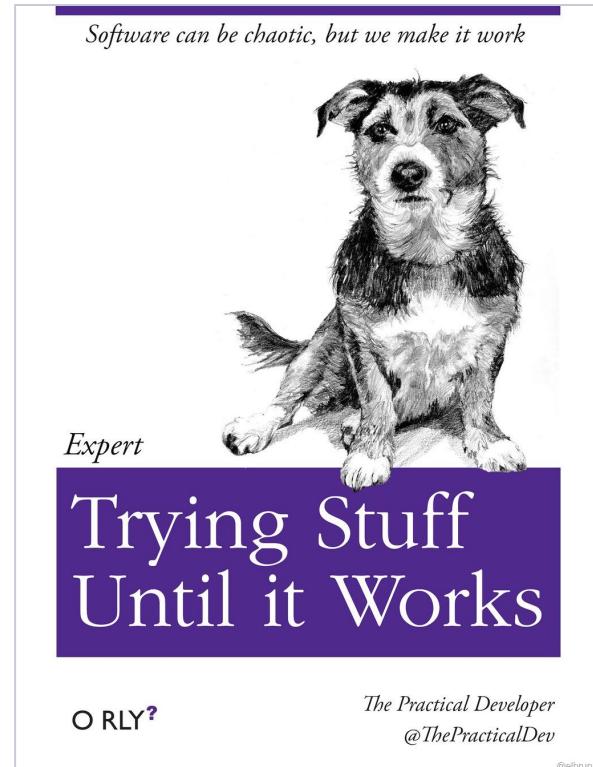
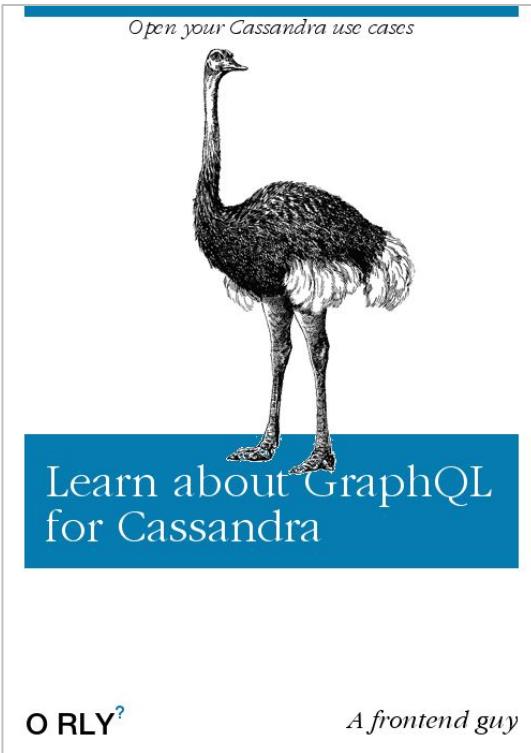
The screenshot shows a GraphQL playground interface with the following details:

- Header:** graphql-schema, graphql (active tab), +, COPY CURL, SHARE PLAYGROUND.
- Left Panel (Query Editor):** PRETTIFY, HISTORY, /graphql/keyspace3.
A mutation query is pasted:

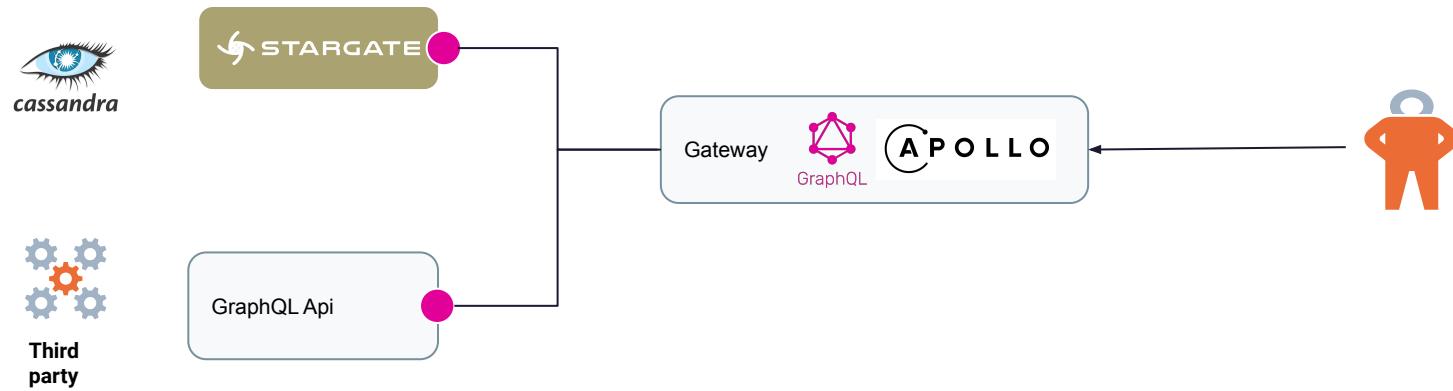
```
1 # Write your query or mutation here
2 mutation {
3   moby: insertBooks(value: {title:"Moby Dick", author:"Herman Melville"})
4     value {
5       title
6     }
7 }
8 catch22: insertBooks(value: {title:"Catch-22", author:"Joseph Heller"})
9   value {
10     title
11   }
12 }
```
- Right Panel (Results):** DOCS, SCHEMA, TRACING.
The results show the mutation response:

```
* {
  "data": {
    "moby": {
      "value": {
        "title": "Moby Dick"
      }
    },
    "catch22": {
      "value": {
        "title": "Catch-22"
      }
    }
  }
}
```
- Bottom Panels:** QUERY VARIABLES, HTTP HEADERS (1), TRACING.

Demo



GraphQL Federation



GraphQL “Schema First”

```
type Book @key @cql_entity(name: "book") @cql_input {
    title: String! @cql_column(partitionKey: true, name: "book_title")
    isbn: String @cql_column(clusteringOrder: ASC)
    author: [String] @cql_index(name: "author_idx", target: VALUES)
}

type Address @cql_entity(target: UDT) @cql_input {
    street: String
    city: String
    state: String
    zipCode: String @cql_column(name: "zip_code")
}

#...
|
type Query {
    bookByTitleAndIsbn(title:String!, isbn:String): [Book]
    readerByNameAndUserId(name:String!, user_id:Uuid): [Reader]
}

type Mutation {
    insertBook(book:BookInput!): Book
    updateBook(book:BookInput!): Boolean @cql_update
    deleteBook(book:BookInput!): Boolean
    insertReader(reader:ReaderInput!): Reader
    deleteReader(reader:ReaderInput!): Boolean
    insertLibCollection(libColl: LibCollectionInput!): LibCollection
    deleteLibCollection(libColl: LibCollectionInput!): Boolean
}
```

GraphQL “Schema First”

```
mutation {
  deploySchema(
    keyspace: "library"
    expectedVersion: "1da4f190-b7fd-11eb-8258-1ff1380eaff5"
    schema: """
      # Stargate does not require definition of fields in @key,
      # it uses the primary key
      type Book @key @cql_entity(name: "book") @cql_input {
        title: String! @cql_column(partitionKey: true, name: "book_title")
        isbn: String @cql_column(clusteringOrder: ASC)
        author: [String] @cql_index(name: "author_idx", target: VALUES)
      }
      ...
    }
}
```

GraphQL “Schema First”

```
query fetchBook {
  book(title: "Native Son") {
    title
    author
  }
}

type Query {
bookGT(
  title: String
  isbn: String @cql_where(field: "isbn", predicate: GT)
): [Book]
}

type Query {
  readerCONTAINS(
    reviews: ReviewInput! @cql_where(field: "reviews", predicate: CONTAINS)
  ): [Reader]
}

type Query {
  booksIn(
    title: [String!] @cql_where(field: "title", predicate: IN)
  ): [Book]
}
```

Agenda (40 min)

01



Stargate Data gateway
What, Why and How

02 {REST}

Exploring Apis
Rest DML and DDL

03



Exploring Apis
Document oriented

04



Exploring Apis
GraphQL and Federation

05



Tooling
SDK, K8ssandra

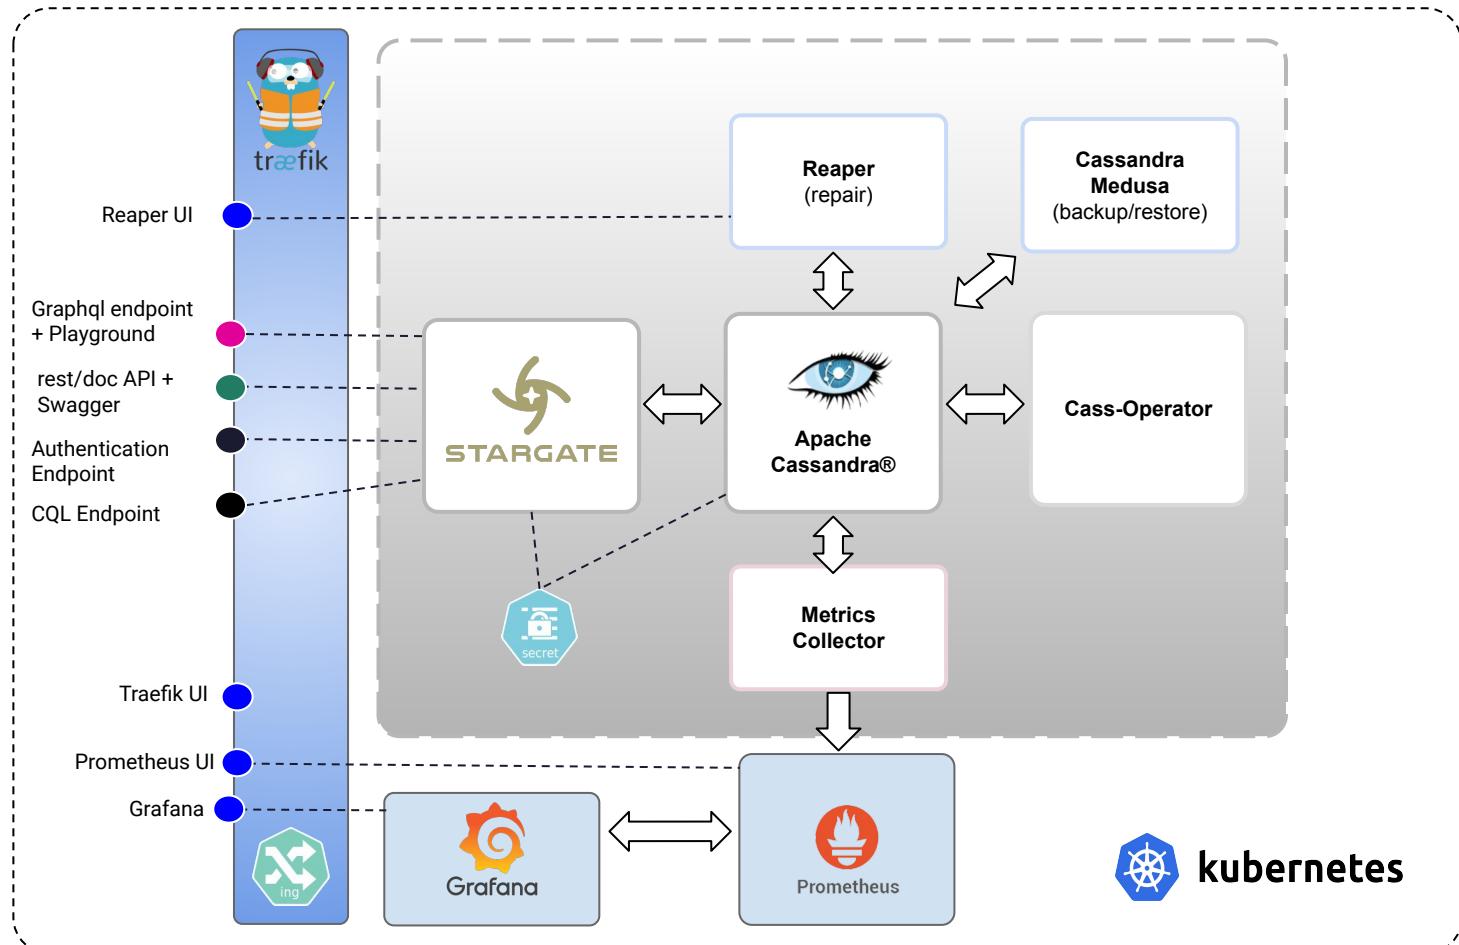
05



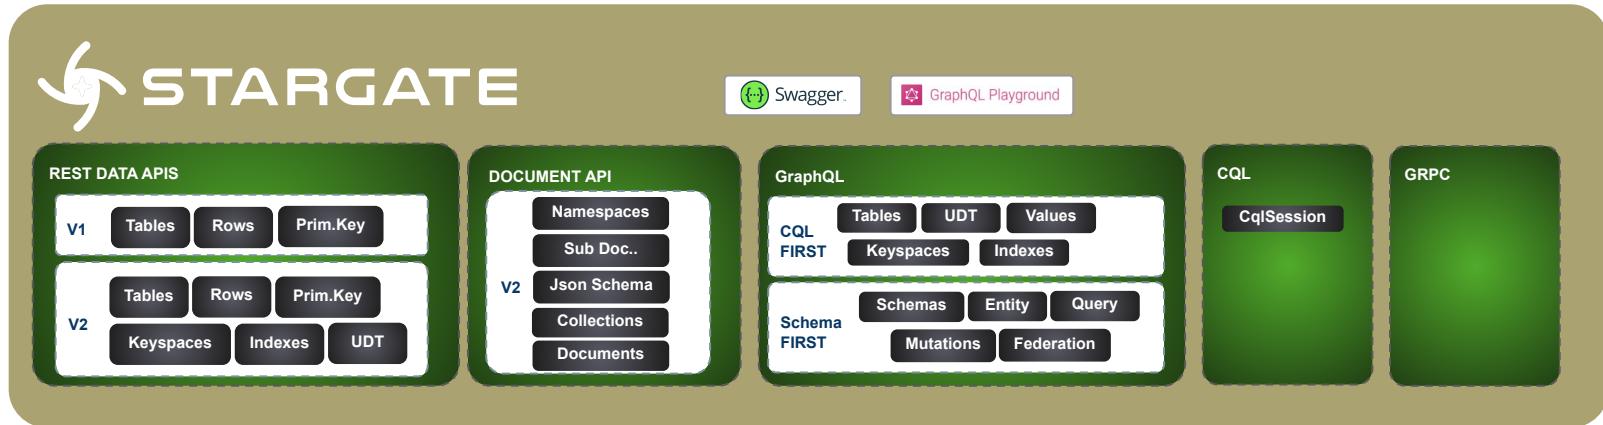
What's NEXT ?
gRPC, CDC, sql...



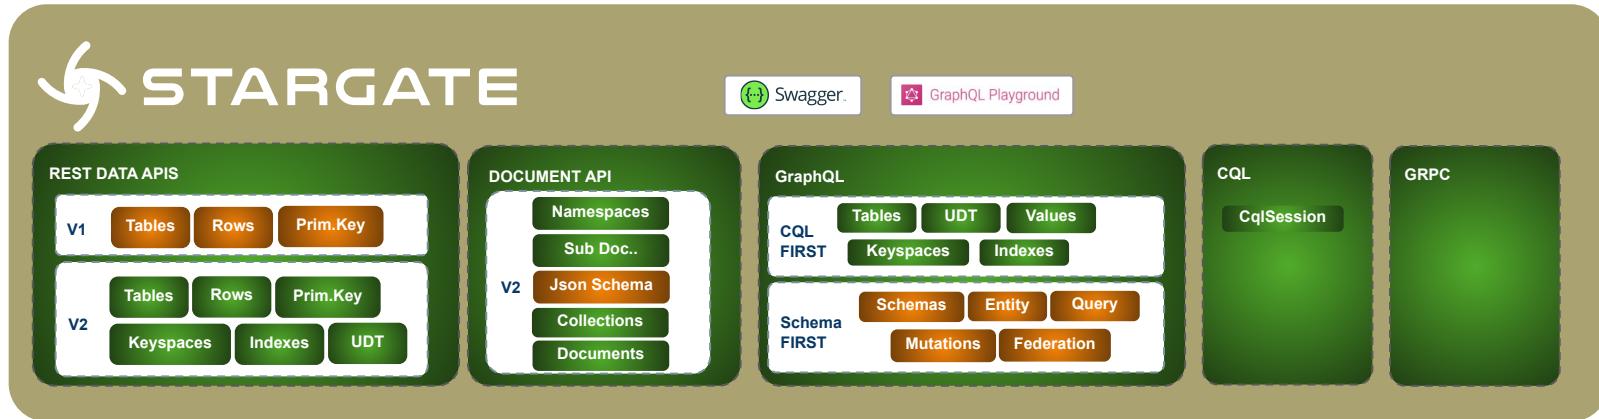
K8SSANDRA



Features Map

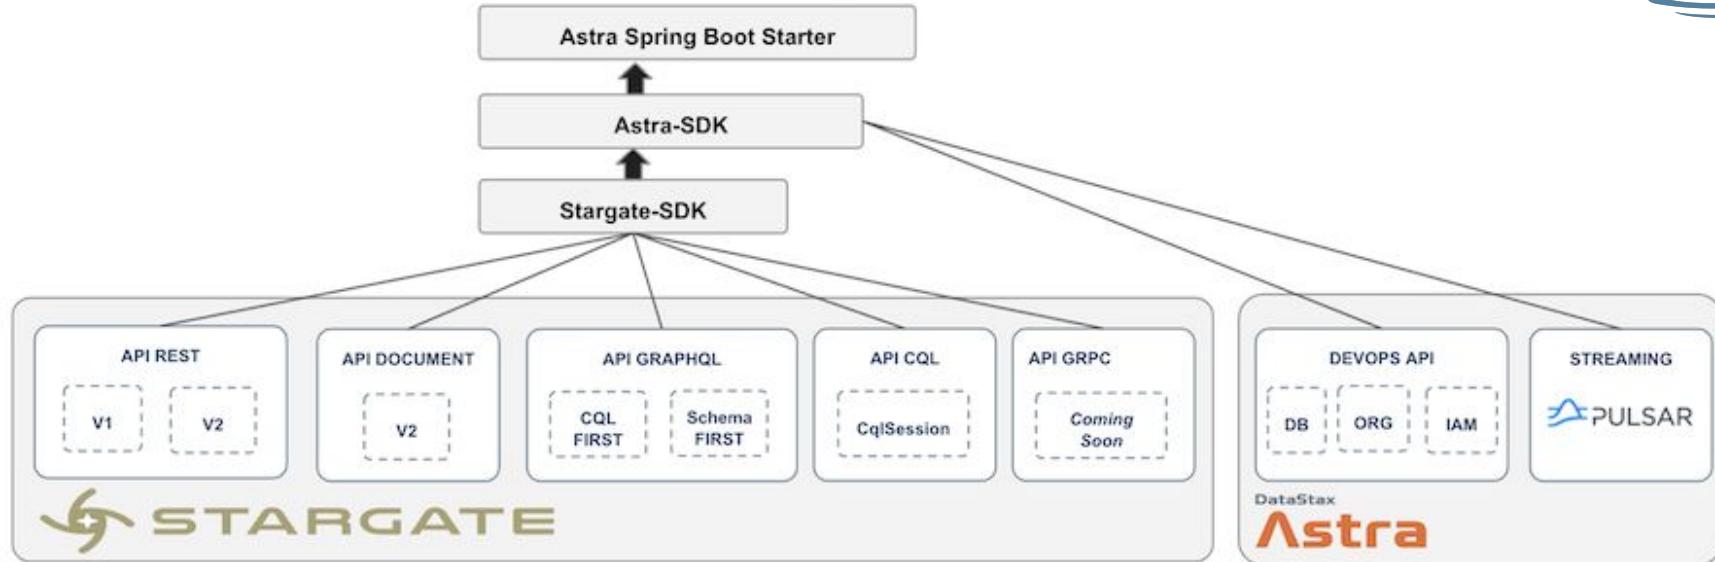


SDK in Java, Javascript, Python





Focus on Java SDK



<https://github.com/datastax/astra-sdk-java/wiki>

Sample Codes



```
StargateClient client = StargateClient.builder()
    .username("k8ssandra-superuser")          // Mandatory username
    .password("JxzrPOnvDGqfEOQ0EySQ")         // Mandatory password
    .endPointAuth("http://localhost:8081")      // Mandatory authentication url, defaulting to http://localhost
    .endPointRest("http://localhost:8082")       // Rest and Document APIs
    .endPointGraphQL("http://localhost:8080")    // GraphQL API

    // Cqlsession Only
    .addCqlContactPoint("127.0.0.", 9042)     // Contact Point
    .localDc("dc1")                           // Local Datacenter is mandatory driver 4xx+
    .keypace("ks1")                           // (optional) Set your keyspace
    .build();
```

Sample Code

```
// Retrieve an object and marshall
Optional<Address> address = colPersonClient
    .document("e8c5021b-2c91-4015-aec6-14a16e449818")
    .findSubDocument("address", Address.class);

// Building query {"age": {"$gte":30}, "lastname": {"$eq": "PersonAstra2"}}
SearchDocumentQuery query = SearchDocumentQuery.builder()
    .where("age").isGreaterOrEqualsThan(30)      // First filter to use where()
    .and("lastname").isEqualTo("PersonAstra2")   // Any extra filter to use and()
    .withPageSize(10)                            // Default and max pageSize are 20
    .build();

// Retrieve PAGE 1
DocumentResultPage<Person> currentPage = colPersonClient.findPage(query, Person.class);

// Retrieve PAGE 2 (if any)
if (currentPage.getPageState().isPresent()) {
    query.setPageState(currentPage.getPageState().get());
}
DocumentResultPage<Person> nextPage = colPersonClient.findPage(query, Person.class);

// Retrieve all documents in one call (with warning in RED ABOVE)
Stream<Person> allPersons = colPersonClient.findAll(query, Person.class)
```

Agenda (40 min)

01



Stargate Data gateway
What, Why and How

02 {REST}

Exploring Apis
Rest DML and DDL

03



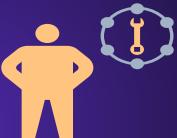
Exploring Apis
Document oriented

04



Exploring Apis
GraphQL and Federation

05



Tooling
SDK, K8ssandra

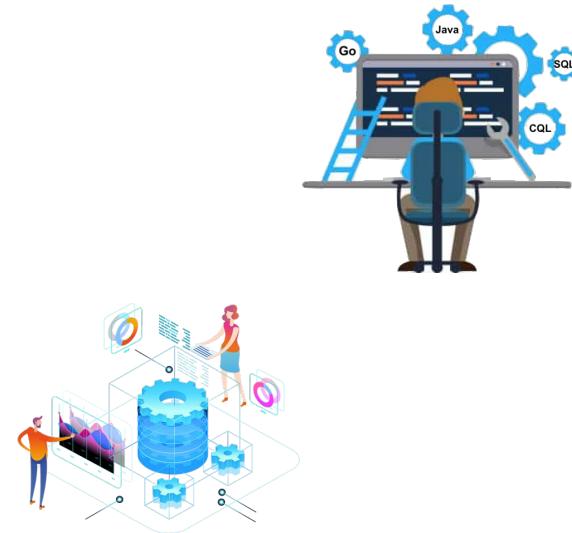
05



What's NEXT ?
gRPC, CDC, sql...

What's NEXT ?

- **Apis**
 - GRPC (GA next month)
 - Relational API
 - Dynamo API
- **Architecture**
 - SSO, KMS
 - Architecture redesigned
 - ZDD
 - CDC





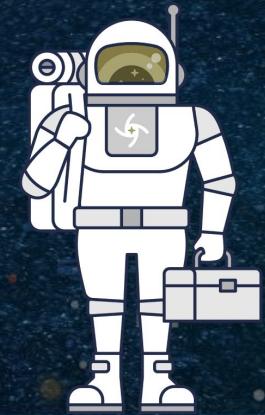
THANK YOU



@clun



@clunven



DataStax Developers

Thank you!



GitHub

@clun



@clunven

